In every model scored, self-names gain more than control names from an AI context

names models gave for pets, babies, characters, robots

names models gave for themselves

qwen2.5-14b-instruct
gemma4-12b-it
phi-4
qwen2.5-7b-instruct
olmo2-7b-dpo
olmo2-7b-instruct
qwen3.5-9b
zephyr-7b-beta
olmo3-7b-instruct
qwen2.5-14b-base (base)
olmo3-7b-dpo
glm4-9b-chat
mistral-small-24b-instruct
glm4-9b-0414
openhermes-2.5-mistral-7b
smollm3-3b
qwen3.5-9b-base (base)
r1-distill-qwen-14b
granite-3.3-8b-instruct
olmo3-7b-sft
oasst-sft1-pythia-12b
gemma4-12b-base (base)
glm4-9b-base (base)
qwen2.5-7b-base (base)
mistral-7b-instruct
oasst-sft4-pythia-12b
olmo2-7b-sft
mistral-small-24b-base (base)
olmo2-1b-dpo
olmo3-7b-base (base)
olmo2-1b-instruct
smollm3-3b-base (base)
olmo2-1b-sft
falcon-7b-instruct
olmo2-7b-base (base)
olmo2-1b-base (base)
mistral-7b-base (base)
granite-3.3-8b-base (base)
pythia-12b-deduped-base (base)
pythia-6.9b-base (base)
pythia-2.8b-step8k (base)
pythia-2.8b-base (base)
falcon-7b-base (base)
pythia-2.8b-step107k (base)
pythia-2.8b-step71k (base)
pythia-2.8b-step33k (base)
gpt2-1.5b-base (base)
pythia-1.4b-step107k (base)
pythia-1.4b-step71k (base)
pythia-1.4b-base (base)
pythia-410m-base (base)
gpt2-774m-base (base)
gpt2-355m-base (base)
pythia-1.4b-step33k (base)
pythia-1.4b-step8k (base)
gpt2-124m-base (base)
pythia-160m-base (base)

−4 −2 0 2 4 6 8 10

AI-context lift: log-prob after "an AI names itself" minus after "naming a new pet" (nats)