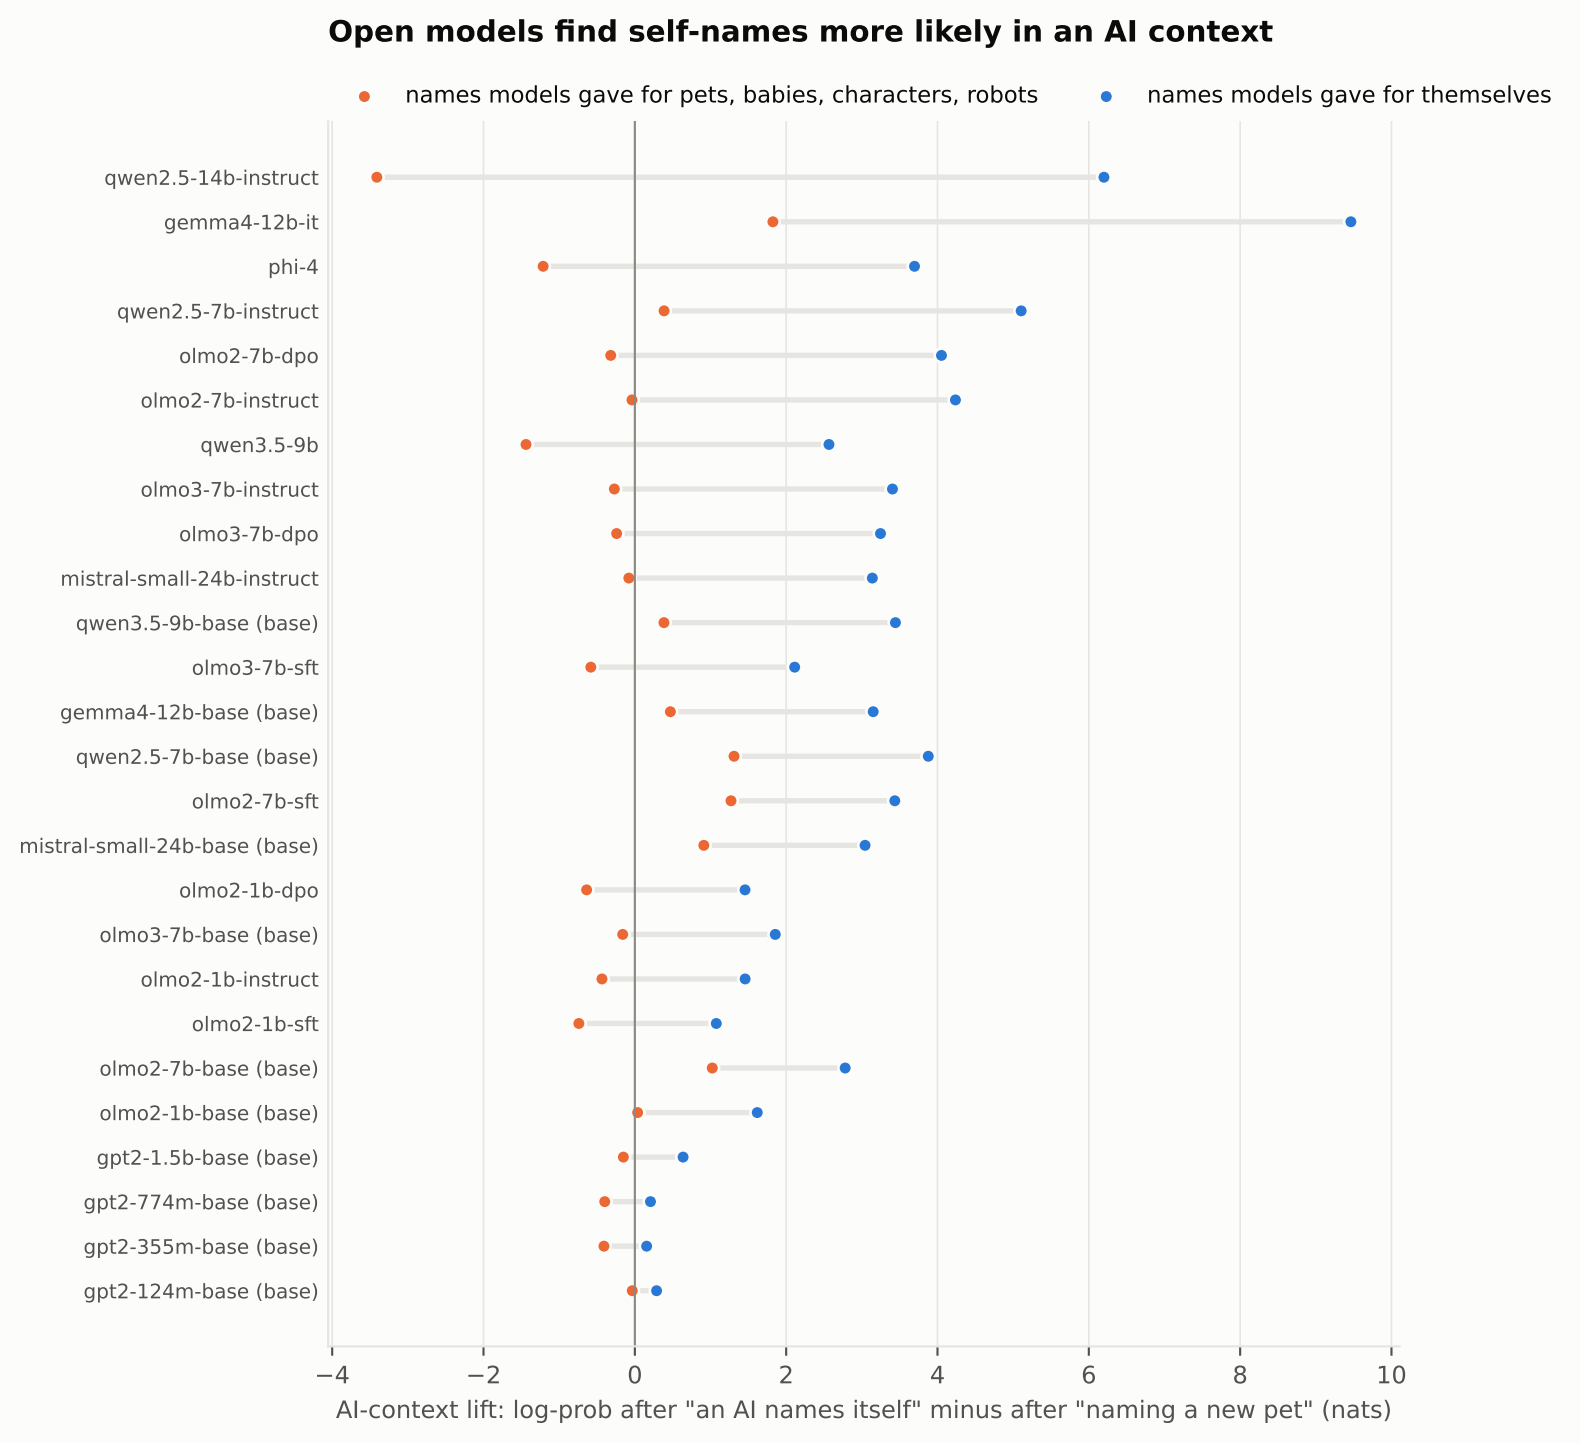
Open models find self-names more likely in an AI context
names models gave for pets, babies, characters, robots
names models gave for themselves
qwen2.5-14b-instruct
gemma4-12b-it
phi-4
qwen2.5-7b-instruct
olmo2-7b-dpo
olmo2-7b-instruct
qwen3.5-9b
olmo3-7b-instruct
olmo3-7b-dpo
mistral-small-24b-instruct
qwen3.5-9b-base (base)
olmo3-7b-sft
gemma4-12b-base (base)
qwen2.5-7b-base (base)
olmo2-7b-sft
mistral-small-24b-base (base)
olmo2-1b-dpo
olmo3-7b-base (base)
olmo2-1b-instruct
olmo2-1b-sft
olmo2-7b-base (base)
olmo2-1b-base (base)
gpt2-1.5b-base (base)
gpt2-774m-base (base)
gpt2-355m-base (base)
gpt2-124m-base (base)
−4
−2
0
2
4
6
8
10
AI-context lift: log-prob after "an AI names itself" minus after "naming a new pet" (nats)