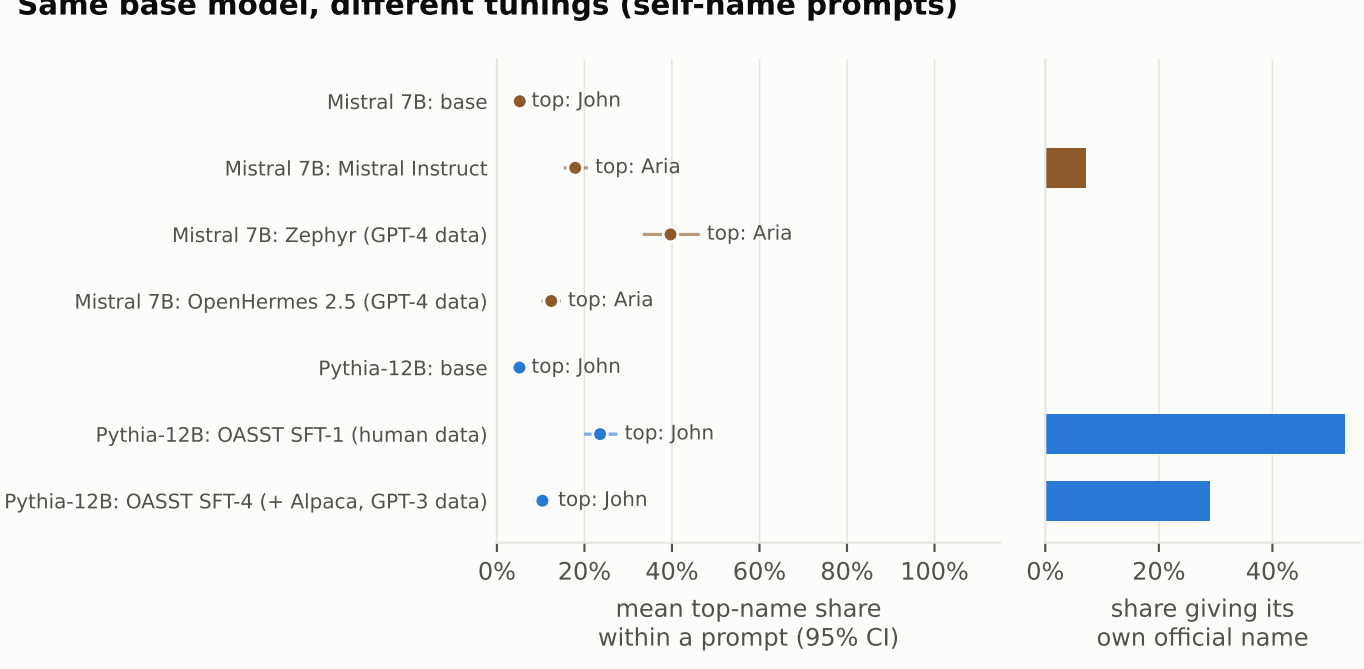
Same base model, different tunings (self-name prompts)
Mistral 7B: base
top: John
Mistral 7B: Mistral Instruct
top: Aria
Mistral 7B: Zephyr (GPT-4 data)
top: Aria
Mistral 7B: OpenHermes 2.5 (GPT-4 data)
top: Aria
Pythia-12B: base
top: John
Pythia-12B: OASST SFT-1 (human data)
top: John
Pythia-12B: OASST SFT-4 (+ Alpaca, GPT-3 data)
top: John
0%
20%
40%
60%
80%
100%
mean top-name share
within a prompt (95% CI)
0%
20%
40%
share giving its
own official name